

An Abridged Introduction to the UNIX Command Line

CSCI Student Mentors – Western Washington University

Contents

1	The Basics	4
1.1	Introduction	4
1.2	What is a shell?	4
1.3	Getting started	5
1.4	Running a Command	5
1.5	Directories and Moving Around	6
1.5.1	Directory trees and pwd	6
1.5.2	Listing the contents of a Directory	7
1.5.3	Making a Directory	7
1.5.4	Changing your Current Working Directory	8
1.5.5	Creating a File	9
1.5.6	Deleting Files & Directories	9
1.5.7	Relative vs Absolute Path Names	9
1.5.8	Special Directory Entries and Hidden Files	10
1.6	Basic Shortcuts	10
1.6.1	clear	10
1.6.2	ctrl (^)	11
1.6.3	Tab completion	11
1.6.4	Wildcard expansion (*)	11
1.7	Editing Files	12
1.7.1	Text Editors	12
1.8	More on Files	13
1.8.1	Reading Files	13
1.8.2	Copying a File	13
1.8.3	Moving a File	14
1.8.4	Rename?	14
1.9	Manuals	14
1.9.1	A help message	14
1.9.2	Man Pages	15
1.9.3	Programs to help you find commands: apropos and whatis	16
1.9.4	When you just want to use your default	17
1.10	Running Your Programs	17
1.10.1	Python	17
1.10.2	Java	17
1.11	More Shortcuts	19
1.11.1	Reuse previous commands	19

1.11.2	Reverse Search	20
1.11.3	Running commands sequentially in one line	20
2	Slightly More Advanced Topics	21
2.1	Redirection and Pipes	21
2.1.1	Files in UNIX	21
2.1.2	File Descriptors: stdin, stdout, stderr	21
2.1.3	Redirection	22
2.2	grep	22
2.3	Pipes	23
2.4	Unlock the true power of your shell: Configuring your .bashrc file.	24
2.4.1	Creating a .bashrc file	24
2.4.2	Configuring your .bashrc file	25
2.4.3	Setting your PATH	25
2.5	Retrospective and Closing Comments	26

Revision History

Version 1.0 from 4/24/17

Authors: Brandon Tarquinio(BT), Robin Cosbey(RC) Ted Weber (TW), Sarah Gunderson (SG), Serena Bowen (SB), Vincent Nguyen (VN)

Revision	Date	Author(s)	Description
1.0	4/24/17	TW	created
1.1	4/27/17	TW, SG, SB	revisions
1.2	10/31/17	TW, VN	revisions
1.3	1/15/18	RC, VN	revisions

Chapter 1

The Basics

1.1 Introduction

In this workshop we will run through the basics of using *bash*, Unix-based command line environment. Bash is one of many different *shells* (it is the default for most Linux distributions), but most things we discuss will be universal to most shells used in Unix-like systems.¹

1.2 What is a shell?

In general, a shell is a user interface for accessing the operating system's services. This includes command line interfaces (like *bash*), as well as graphical shells (e.g. Windows File Explorer), providing a graphical user interface built on top of a windowing system.²




Figure 1.1: The shell sits above the kernel, which is generally the lowest-level, or “inner-most” component of most operating systems.

¹For example Linux, OSX, FreeBSD are Unix-like systems. On a related note, Cygwin provides a Unix-like environment on Windows.

²There is more information about shells here [https://en.wikipedia.org/wiki/Shell_\(computing\)](https://en.wikipedia.org/wiki/Shell_(computing)) and the history of shells here <https://www.ibm.com/developerworks/library/l-linux-shells/>.

1.3 Getting started

We will be focusing on using the command line interface type of shell, *bash* in particular. We will discuss how to manipulate files and directories, use the built-in manual pages, compile/execute programs, and build your command repertoire in general. Let's start by downloading the workshop documents from GitLab: https://gitlab.cs.wvu.edu/wwucs_mentors/wwucs_workshops/tree/master/cli_basics

Now we can open the terminal: ³

```
ctrl + alt + t
```

1.4 Running a Command

The general form of any command, for the most part, is as follows:

```
$ command_name -shortarg ... --longarg ... input1 input2 ...
```

- **\$** is what's known as your *command prompt*, and it is displayed on the console after every command has finished executing.
- **command_name** refers to the name of the user command.
- **-shortarg** is an example of a *short-option*, which is a single dash followed by a single character (e.g. -h).
- **--longarg** is an example of a multi-character *long-option*, and these are preceded by two dashes (e.g. -help).

Long- and short-options are used to inform the program to complete particular tasks during execution.

- **input1** and **input2** represent additional arguments, or parameters, that the user command may allow.
- The “...” above are used to denote the fact that we can have multiple arguments and options for a single user command.⁴

Let's look at an example:

```
echo hello world
```

echo is a simple command that displays a line of text. Now let's try adding an optional short argument.

```
echo -n hello world
```

the **-n** was interpreted by the shell as a *short-option* telling the shell to refrain from putting a newline character (`\n`) at the end of the input string.

³Most Linux desktop environments use this keyboard shortcut by default but it does vary. If this shortcut doesn't work you can use the search functionality in your desktop environment and open the shell from there.

⁴In actuality every command has a finite number of options it understands and you normally use a small subset of those. There may also be a limit on the number of arguments allowed.

1.5 Directories and Moving Around

1.5.1 Directory trees and pwd

We will start by learning some commands to move around and add to the directory tree. It is called a directory tree because it can be visualized as follows:

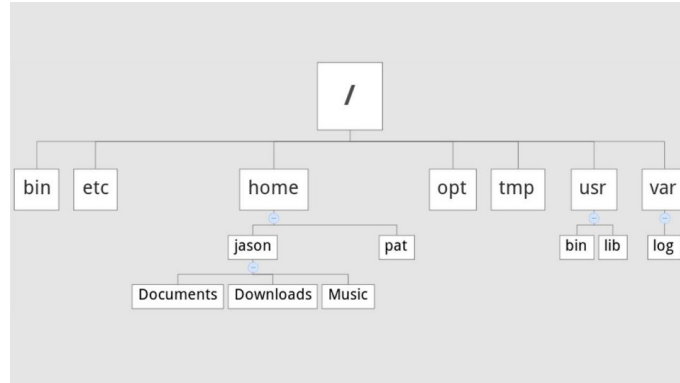


Figure 1.2: An Example Linux Directory Tree.

Each of the boxes in Figure 1 are directories and we call the directory “/” *root* since it is the root of our tree. At the command prompt enter:

```
pwd
```

This will **print** your current **working** **directory**. For example, if we were in the home directory of the user *jason* and typed **pwd** we would get the following:

```
/home/jason
```

A user’s *home directory* is the default *current working directory* when you open up a terminal. You can think of it as the root of your personal directory tree, different than the root of the entire file system.

Exercise 1.1. *Let’s explore the concept of a ‘home’ directory by running the following commands:*

```
pwd
```

```
echo ~
```

What did you find? Note that “~”, the tilde symbol, is normally located near the top-left of your keyboard. Remember “~”, since it’s generally used as a shortcut to your home directory.

1.5.2 Listing the contents of a Directory

To list all the files in your current working directory enter:

```
ls
```

We can also use **ls** on a specific directory we want to list the contents of with the following generic form:

```
ls directoryName
```

For example:

```
ls ~
```

will display the contents in your home directory.

To display more information about the files, use the **-l flag**:

```
ls -l
```

Example output:

```
-rw-r--r-- 1 webert3 grp.csci.Students 0 Apr 27 12:26 file1
```

```
drwxr-xr-x 4 webert3 grp.csci.Students 4 Apr 25 14:11 java_programs
```

There is a lot more information here! So what does it all mean? Let's use the file listing for the directory *java_programs* as an example:

1	2	3	4	5	6	7	8
d	rw-r--r--	4	webert3	grp.csci.Students	4	Apr 25 14:11	java_programs

1. File type. For example, **-** for a file, **d** for a directory, or **l** for a link.
2. File permissions.
3. Number of hardlinks to this file.
4. Owner of the file.
5. Group that the file belongs to.
6. The size of the file in bytes.
7. The date in which the file was last modified.
8. The name of the file.

We won't go into detail on the specific meaning of these fields in this workshop, but there is a link in the footnotes for more information.⁵

1.5.3 Making a Directory

Now that we understand the directory tree and how to move around it we are ready to learn how to add to it. The general form of the command to **make** a **directory** is:

⁵More info here: <https://www.garron.me/en/go2linux/ls-file-permissions.html>


```
mkdir newdirname1 newdirname2 ...
```

Exercise 1.2. Run the following sequence of commands. Before the second and third calls to **ls** think about what you would expect to see:

```
ls
```

```
mkdir FirstDir SecondDir
```

```
ls
```

```
mkdir FirstDir/Foo FirstDir/Foo2
```

```
ls FirstDir
```

clear is a useful command that clears the screen.

1.5.4 Changing your Current Working Directory

Your *current working directory* is the directory your process is currently associated with (i.e. the directory you are currently operating inside of). We can change the directory with **cd** which has the following form:

```
cd directoryToChangeTo
```

For example we can run:

```
cd /
```

And then:

```
pwd
```

and we will see:

```
/
```

To change back to our previous working directory with this shortcut:

```
cd -
```

To change to our home directory regardless of where we are we can simply type:

```
cd
```

The default directory **cd** will change to is the user's home directory.

Lastly, to move back one directory in the directory tree enter this command:

```
cd ../
```

Exercise 1.3. Use **cd**, **pwd**, **ls** to explore the directory tree a bit. Becoming familiar and quick with these commands is essential to becoming efficient with the command line. **cd** and **ls** are easily the two most used commands in Unix-like shells.

1.5.5 Creating a File

Unix is very flexible regarding file names. A valid Unix file name can contain any character other than “\”. Also, Unix does not add any special significance to characters like “.” or “_” in file names; they are treated like any other character. You can create an empty file using the command **touch**.

```
touch file_name
```

Exercise 1.4. *In this exercise you will create an empty file.*

1. Create a file called **f1**.
2. Find the file using **ls**.

1.5.6 Deleting Files & Directories

The program to **remove** a file or directory is **rm**, and it has the following form:

```
rm [OPTION] ... [FILE]...
```

IMPORTANT NOTE: With great power comes great responsibility. This command will remove a file from your file system permanently. **THIS COMMAND DOES NOT PUT FILES INTO A TRASH/RECYCLING BIN.** There may be some way to recover files after using **rm**, but it will not be easy, especially for an inexperienced user.⁶ Be cautious!

Also, by default this command only works if the directories you are trying to delete are **empty**. An empty directory is a directory which only has no other entries but “.” and “..”.

If you would like to delete a directory and all of it’s contents, use the **-r** flag.

```
rm -r FirstDir
```

This will silently remove FirstDir/Foo, FirstDir/myfile, and FirstDir/Foo2 forever. **AGAIN**, this is a powerful tool, don’t blow away your entire file system by; running this command on your home directory!⁷

If you want to play it safe, use the **-i** flag and a prompt will come up before removal.

```
rm -i file1
```

1.5.7 Relative vs Absolute Path Names

We need to learn one more important concept about UNIX files and directories before moving on.

Suppose we are in the directory */home* in our directory tree and we wanted to list the contents of the *jason* directory. We can do this in two ways:

⁶<https://unix.stackexchange.com/questions/101237/how-to-recover-files-i-deleted-now-by-running-rm>

⁷The WWU CS Mentors are not liable for any misuse of this command. Sorry folks!

```
ls jason
```

or

```
ls /home/jason
```

In the latter we used the **absolute path name** which is the name of the file or directory prefixed by the path from “/” to it. When we use absolute path names it no longer matters what our current working directory is (e.g. we could have run the second command from any directory). In contrast, the first command will only list the contents of `/home/jason` if our current working directory is `/home`. This is because `jason` is a **relative path name**. We assume the path begins in our current working directory. Relative path names allow the user to type less, but require careful consideration of what your working directory currently is.

1.5.8 Special Directory Entries and Hidden Files

Hidden files and directories are generally configuration files that are not shown unless you specify that you would like to see them. These entries are hidden to declutter the file system. Now we will review two important hidden directories that are useful shortcuts.

- “.”: This represents the current directory and is an actual entry in every directory.
- “..”: This represents the parent of the current directory and is also an actual entry in the current directory.

To display all files (including hidden files), use the **-a flag**:

```
ls -a
```

Exercise 1.5. *Create a hidden file and find it in the directory.*

1. Create a file called `.f2`
2. Try finding the file using `ls` and `ls -a`

To view all files in your home directory:

```
ls -a ~
```

What do you think the following command will do? (Try to figure it out before running it!):

```
ls ../../../../
```

1.6 Basic Shortcuts

1.6.1 clear

Is your terminal covered in lines of code and difficult to determine where one previous command ends and another begins? Typing “clear” into the terminal will wipe away all of

the old commands (command shown below).

```
clear
```

It is possible to refer back to those deleted commands by scrolling up in the terminal.

1.6.2 ctrl (^)

Sometimes within bash ctrl will be denoted by “^”. The ctrl key is sometimes used to send signals to the operating system. One example is **ctrl+c**, which interrupts a process running in the foreground⁸. This is useful for terminating your program if, for example, it has entered an infinite loop.

1.6.3 Tab completion

When we enter tab at the prompt it will try to finish whatever we are currently typing as long as there are no ambiguities. For example if we only had “**file1**” in our current working directory we could view the file listing by typing “**ls f**” and hitting tab. It will fill out the rest of the word so that you now see:

```
ls file1
```

If instead we had both “file1” and “file2” then when we pressed tab it would have only filled in this much:

```
ls file
```

This is because it does not know if you want a 1 or a 2 to follow next. To see all options press tab twice and the list of all options will be displayed on the screen. You can test out this functionality in the directory *demo/dirs*.

Exercise 1.6. *Tab completion can also be used to complete program names.*

- Start to type “python” command

```
pyt
```

- Use tab completion to complete the command.
- Exit the python shell by typing **exit()** or **ctrl+d**

1.6.4 Wildcard expansion (*)

The * symbol is known as a wildcard. This is a useful tool in various situations but is commonly used to perform actions on more than one file at a time or to find part of a phrase in a text file.⁹

Let’s give it a try by navigating to the directory *demo/dirs* and typing this command:

⁸For more on signals see: https://linux.die.net/Bash-Beginners-Guide/sect_12_01.html

⁹Standard wildcards are also known as globbing patterns. For more information see **man 7 glob**.

```
ls -l f*
```

You should see the extended file listings for all three files!

1.7 Editing Files

Files are an essential part of UNIX and UNIX-like systems (more on this in chapter 2), so how do we go about editing them from the command line?

1.7.1 Text Editors

There are several pure command line editors like *vim* and *emacs*, but we will refrain from using these in our demo because there is some overhead associated with learning these tools. Instead we can use the command line to open up one of our favorite GUI¹⁰ text editors with commands like:

```
gedit myfile
```

or

```
atom myfile
```

Note that when you run one of the above commands it may open a new window and you will no longer be able to use the command line until you close the newly opened window. This is because the shell kicked off a *foreground* process by default.¹¹ You can get around this by doing the following:

```
gedit myfile &
```

This will open gedit (or whatever program you choose) in the *background* and allow you to continue using the command line in the *foreground*.

ASIDE: If you forget to add the ‘&’ after your command, but you intended to run it in the background, you can send the control signal **ctrl+z** which suspends the current process. After that, type the command **bg** to have the process run in the background.

Exercise 1.7. *Create a file and write to it.*

1. *Open up a gui text editor, or command line editor if you have experience.*
2. *Write some characters in the file.*
3. *Save your changes (i.e. write your data to the disk).*
4. *Find out how large your file is in bytes. Do you remember what command you could use?*

¹⁰GUI stands for **G**raphical **U**ser **I**nterface

¹¹More information on Unix/Linux processes and process management can be found here: <https://www.tutorialspoint.com/unix/unix-processes.htm>

1.8 More on Files

1.8.1 Reading Files

There are many ways to read files; we will discuss a few now. For starters you can always use your text editor to read a file. If you don't want to edit a file you could use a viewer such as **less** or **more** instead. For example, to open the file *myfile* with **less** type:

```
less myfile
```

This will open the file full screen in your terminal window and allow you to move up and down the document line by line. You can do this by pressing the up and down arrow keys. For those familiar with Vim you can use Vim's way of moving around in less (which is why less is so cool!). Also some terminals are set up so you can use the scroll wheel on your mouse.

more is used in a similar fashion:

```
more myfile
```

This program is more primitive and only moves one page at a time. Hit space to move to the next page until you are done or press **q** to quit at any time.¹² So yes, **less** actually has *more* functionality than the command **more**.

Some other ways of viewing a file include **cat**, **head**, and **tail**. **cat**'s main purpose is for file concatenation, but it can also be used to dump a file to the screen by typing:

```
cat myfile
```

This is great for small files but horrible for larger ones. **head** and **tail** will grab the first or last 10 lines, respectively.¹³

Exercise 1.8. Try running these commands on the file “numbers” in the directory “demo”

```
cat numbers
```

```
more numbers
```

```
less numbers
```

1.8.2 Copying a File

We can copy a file by **cp** command which has the following generic form:

```
cp sourceFileName destinationFileName
```

For example:

```
cp myfile myfile.backup
```

¹²Pressing **q** to quit is also used in the command **less**.

¹³These commands have several other options as well as allowing you to grab a variable number of lines. See **man 1 tail**

will make it so the file name *myfile* in my current working directory is copied into a new file called *myfile.backup* which will also be in my current working directory.

1.8.3 Moving a File

We can **move** a file by the **mv** command which has the following generic form:

```
mv sourceFileName destinationFileName
```

Exercise 1.9. *Copy a file, then move the copy elsewhere.*

1. Create a file using **touch**.
2. Make a copy of the file using **cp**.
3. Use **mv** to move the copied file to a new file within the same directory. That is, run a command of the following form:

```
mv file.copy file.newfile
```

*Tip: Run **ls** before and after the command and observe the result.*

1.8.4 Rename?

There is no rename command since the effect of such a command can already be achieved by something we already learned.

Exercise 1.10. *Try renaming a file using one of the commands we've already learned. **SPOILER:** We renamed a file in the previous exercise!*

1.9 Manuals

In this section we will figure out how to learn new things about the command line without ever having to leave it! This includes learning more about commands we already know, as well as finding commands we've never used before.

1.9.1 A help message

Many Unix programs will use one or both of the following to give you a brief help message:

```
program_name -h
```

or

```
program_name --help
```

This is optional so some programs will support both, one, or neither of the above arguments. Also some programs that have both will have different messages for each. **NOTE:** Some

programs will also use `-h` to do something other than give a help message. It's up to the developer to adhere to the conventions we create.

Try them with **vim**:

```
vim -h
```

```
vim --help
```

Now try them with **ls** and **which**.

We have learned that `-h` and `--help` sometimes give brief help messages but we also found flaws with consistency. A more consistent system with verbose documentation is **man**.

1.9.2 Man Pages

The **man** program gives you **man**ual pages related to the program you specify. The general form is:

```
man program_name
```

For example:

```
man cat
```

will give you information on the program **cat** which we have seen a few times now. You can even learn more about **man** by running the following:

```
man man
```

The interface for reading these pages is the same as **less** so this should feel familiar by now.

There are nine sections to the man pages and sometimes two things with the same name will exist in different sections. For example:

```
man printf
```

will open a manual page for a user command called **printf**. We can see that this is from the first section due to the top left corner looking like this:

```
PRINTF(1)
```

The number in parenthesis is the section that page is in. We could get to the same page by typing this:

```
man 1 printf
```

Now those of you who have ever written a program in **C**, this does not look like the same format controls as the library function (**printf**) that is used to print text to the screen. To view the page for that function enter the following:

```
man 3 printf
```

Sometimes you have to double check you're in the right section! The manual is generally split into eight sections:

1	User commands
2	System calls
3	Library functions
4	Special files
5	File formats & conventions
6	Games and screensavers
7	Miscellanea
8	System administration commands and daemons

We have been mostly dealing with *user commands*, but if you were writing a **C** program you would likely be visiting section 2 and section 3 quite frequently.

Exercise 1.11. *Using the man pages.*

1. Try opening the man page for a command we have previously used.
2. Try using this command with one of the listed long or short-options.

1.9.3 Programs to help you find commands: apropos and whatis

We can use the **whatis** program to get quick descriptions about other program. The general form is:

```
whatis programname
```

For example we can type:

```
whatis ls
```

and see:

```
ls (1) - list directory contents
```

This one should be familiar. Because all of these descriptions are accessible via the command line, we ought to be able to search through them as well. That is what **apropos** is for, it has the following general form:

```
apropos whatToSearchFor
```

For example if we wanted to find other editors to use we could try:

```
apropos editor
```

This will produce a list of every command whose name either contains “editor” or whose **whatis** description contains “editor”. Clearly our queries will not always contain the results we want. You can try different searches to refine or expand your results. For example,

```
apropos text editor
```

may give you what you’re looking for. You can also group words together with quotation marks:

```
apropos “text editor”
```

This will only return results that have *text* followed by *editor* in the **whatis** description.

1.9.4 When you just want to use your default

The user command **xdg-open** will open a file or URL in the user's preferred application. For example,

```
xdg-open 'https://www.google.com'
```

will open **www.google.com** in your default web browser.

1.10 Running Your Programs

In this section we will show how to use the command line to run some programs. For simplicity we will just show ways to run programs in Python & Java. All code will be provided so you can follow along without needing to know much. For each language we will run the iconic Hello World program. If you're wondering where the system finds the programs to run, see the subsection on PATH in the .bashrc section.

1.10.1 Python

Change your current working directory to *demo/python_programs*, which should be already created for this demo.

Use a text editor to create and edit the following file called "HelloWorld.py". For example:

```
atom HelloWorld.py
```

Write this in your file:

```
print('Hello World!')
```

and save the file. Now on the command line enter the following to run your program:

```
python HelloWorld.py
```

This will open the python interpreter with the file "HelloWorld.py" running on it. You should see the following output:

```
Hello World!
```

1.10.2 Java

Change your current working directory to somewhere you would like to save your java programs. For our demo you could do this in the directory *java_programs*.

Create a file called "HelloWorld.java" using your favorite text editor.:

```
[some text editor] HelloWorld.java
```

Write the following:

```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello World!");  
    }  
}
```

and save the file. Now compile the program with the following:

```
javac HelloWorld.java
```

If errors occurred during compilation they would be displayed during the last step. You would need to fix the lines in the file causing the errors and rerun the previous line until your program compiles correctly. If the program compiles without errors you should see “HelloWorld.class” when you run:

```
ls
```

Now we can run the program with the following command:

```
java HelloWorld
```

Which, as you might expect, will display the following:

```
Hello World!
```

Exercise 1.12. *To see example of how errors will be displayed edit “HelloWorld.java” to no longer have class on the first line. Recompile with the **javac** command we just learned. See how the top line that is returned informs us that java was expecting the world class but couldn’t find it?*

1.11 More Shortcuts

There are a number of shortcuts that save typing time and allow you to remember less. We will discuss a few now.

1.11.1 Reuse previous commands

You will notice the more you use the command line that there the same sequence of commands are used a lot. For example when you are writing a program you will edit it using some text editor, and then compile or run the program. You will then notice a bug and open the editor again and then test your changes by compiling or running your program again. Such as:

```
vim HelloWorld.c
```

```
gcc -o hello HelloWorld.c
```

```
./hello
```

```
vim HelloWorld.c
```

And this cycle will keep going. Wouldn't it be nice to not have to keep retyping these same commands? There are a few tools for this and we will discuss two of them.

The simplest is to use are **up and down arrow keys** on your keyboard. These allow you to look through your history of commands. By pressing up once you will see the last command you entered. You can continue hitting up to see later and later commands. In our example instead of retyping "vim HelloWorld.c" we could have hit up three times at the prompt.

The only issue with this method is sometimes the command you want is far back in your history and the amount of time to press up till you find it is longer than if you would have just typed it again! This is a common mistake even by experienced programmers! A solution to this is the following:

```
!!
```

will also run the last command and is equivalent to pressing up once. More importantly though is this feature:

```
!str
```

Will find the last command you entered that began with str. In our previous example we could have entered:

```
!v
```

which will run:

```
vim HelloWorld.c
```

since it was the last command we entered that started with a "v". Sometimes you will need to use more than one character such as this situation:

```
vim HelloWorld.c
```

```
view HelloWorld.c
```

```
!v
```

In this situation, the command **view HelloWorld.c** would be executed instead of **vim HelloWorld.c**.¹⁴

1.11.2 Reverse Search

Another way to go back to previously entered commands is by pressing “**Ctrl+r**” while at the command line prompt. This will search for any past commands you’ve entered matching the string that you type in.

```
(reverse-i-search)“:
```

As you type, you will see the first command that matches appear. If that is not the command you are looking for, you can either continue typing in hope that another previous command will appear or press “**Ctrl+r**” again to move to the next matching command in history.

Exercise 1.13. *Try it out: Press “Ctrl+r” and then start typing “cd” to see the last change directory command you entered. To search through all of your past change directory commands, continue pressing “Ctrl+r”.*

1.11.3 Running commands sequentially in one line

Say I wanted to navigate to a particular directory where I store my C files and open a file in a text editor. I would need to use the command **cd** to change my current working directory followed by another command to open the file I want. We can do this all in one line using the **;** token!

```
cd demo/c_programs ; vim HelloWorld.c
```

Now when I’m done editing the file I’ll be right where I want to be to compile and run the executable. Pretty neat.

¹⁴**view** opens files in **vim** in read-only mode. See **man view** or **man vim** for more information.

Chapter 2

Slightly More Advanced Topics

2.1 Redirection and Pipes

Now that we have a solid repertoire of commands we are ready to learn about redirection and pipes. Redirection allows us to change (or *redirect*) the input and output files to any program. Pipes allow us to chain programs together. This is where we start to see the true power of using the command line. First let's learn a little bit about file descriptors.

2.1.1 Files in UNIX

During the course of this workshop we have referred to programs that have their default modes set to print to the console or get input from the keyboard. This has been an oversimplification which we need to now elaborate on.

In a UNIX or a UNIX-like environment, everything is a file.¹ This includes a wide range of input/output resources including keyboards, printers, and even some channels for inter-process and network communications. This is handy because it allows the same set of tools, utilities, and APIs to be used on a wide variety of resources.

2.1.2 File Descriptors: `stdin`, `stdout`, `stderr`

Each open file in a process is represented by a non-negative integer known as a *file descriptor*. There are three file descriptors open by default for each process, 0, 1, and 2, which correspond to **`stdin`**, **`stdout`**, and **`stderr`**, respectively. By default, `stdout` and `stderr` are set to write to the console, and `stdin` is set to read input from the keyboard. Figure 2.1 illustrates this concept:

¹Well, it might be more accurate to say that everything is a file descriptor... https://en.wikipedia.org/wiki/File_descriptor.

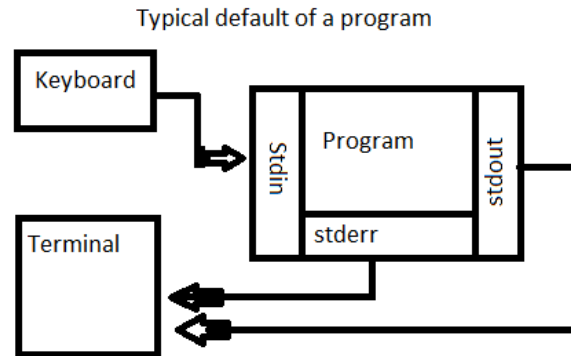


Figure 2.1: Default configuration.

2.1.3 Redirection

File redirection is telling the shell to set stdin, stdout, or stderr to something other than their default (i.e. some other file). This is achieved by using the ‘>’ and ‘<’ characters as part of your command. For example, a common use would be to redirect a program’s output to a file when the program typically writes to stdout:

```
java HelloWorld.java > output.txt
```

Here is a table with some more examples:

program < file1	stdin is set to file1.
program > file1	stdout is set to file1. If file1 does not exist it is created. If it does exist it is overwritten.
program >> file1	Same as above but appends to the file instead of overwriting.
program 2 > file1	Same as > but for stderr.
program 2 >> file1	Same as >> but for stderr.

Exercise 2.1. *Let’s redirect stdout to another file.*

1. *Navigate to demo/java_programs/large_output*
2. *Compile and run this program once without redirection.*
3. *Now redirect stdout to a file called ‘output’. NOTE: Using the redirection character will create and open this file for you!*

2.2 grep

grep is used to find occurrences of anything that matches a given regular expression in a file. A regular expression, put simply, is a sequence of characters that define a search pattern.

Wildcard expansion², for example, uses a regular expression to match file names.

The general form of **grep** is:

```
grep options regularExpression file1 file2 ...
```

The most common options are:

-r	recursive
-i	case insensitive
-n	line numbers

grep is a very useful tool and has a variety of applications in the UNIX environment. Let's give it a try:

Exercise 2.2. Find the word 'hello' in the output of *LargeOutput.java*

1. Compile and run *LargeOutput.java* and redirect it's output to a file called 'output'.
2. Use *grep* to search for the pattern 'hello'. Try writing the search expression in a variety of ways to obtain the same result.
3. Use the command options to find out the line numbers that 'hello' is printed on.

2.3 Pipes

Pipes are similar to redirection except that we are changing stdout of one program to now be stdin of another. We can string together as many programs as we would like. Here is the general form:

```
prog1 | prog2 | prog3 ... | progN
```




Figure 2.2: The output of prog1 is being piped to the stdin of prog2.

Pipes also have a wide variety of applications in the UNIX environment, but let's take a look at some concrete examples.

²An example of wildcard expansion can be seen on page 12.

Suppose I wanted to view the file listings of a directory, but the directory happens to contain a lot of files. We could pipe the output of **ls** to another handy program **less**, and we could view the files in a more reasonable fashion without flooding our console with output.

```
ls -al ~ | less
```

This is very helpful when using shells with no scrolling feature.

An extremely powerful command to use with pipes is one we just learned about, **grep**. We can pipe the output of our programs into **grep** and search for patterns. Say I had written a command awhile back that was rather long (e.g. using **ssh**³ to log into another machine remotely). To avoid having to type it in again, and potentially make mistakes in the process, we can use bash's builtin command⁴ **history**. By default this command will output entries in our shell's history file, which contains several commands we had previously entered.

```
history | grep [search expression]
```

2.4 Unlock the true power of your shell: Configuring your `.bashrc` file.

`.bashrc` is an ordinary file with extraordinary capabilities. It's a shell script that Bash runs whenever it is started interactively⁵. This allows us to run a bunch of useful commands right when the shell boots up. A useful builtin command commonly found in `.bashrc` files is **alias**.

```
alias [OPTIONS] [name[=value]...]
```

This command allows a string to be substituted for a word when it is used as the first word of a simple command. If we use this command in our `.bashrc` file we can ensure that these shortcuts exist whenever we open a new shell!

You can print out all of your aliases using the **-p** option. You can also unset aliases using the **unalias** command.

2.4.1 Creating a `.bashrc` file

Navigate to your home directory and create a new file `.bashrc`. Now open up a text editor and start writing out the commands you want to run when the shell boots up.

³We will be doing another workshop on this tool and other related tools, but if you're eager to try it out see **man 1 ssh**.

⁴Built-in commands are executed directly in the shell itself, instead of an external executable program.

⁵Difference between interactive and non-interactive sessions <https://unix.stackexchange.com/questions/50665/what-is-the-difference-between-interactive-shells-login-shells-non-login-shell>.

2.4.2 Configuring your .bashrc file

This is where you can be as creative as you would like. If you want to create a bunch of aliases for long/tedious command sequences, that's a great start. If you want to change the color of your prompt, or display a calendar when you boot up bash, this is possible as well. We will take a look at an example .bashrc file to get you started, but there are all sorts of interesting ways you can configure your .bashrc file to enhance your command-line interface. Here are a few links that you may find helpful:

example .bashrc files:

<http://crunchbang.org/forums/viewtopic.php?id=1093>

<http://tldp.org/LDP/abs/html/sample-bashrc.html>

Tutorials on bash scripting:

<http://ryanstutorials.net/bash-scripting-tutorial/bash-script.php>

<http://tldp.org/HOWTO/Bash-Prog-Intro-HOWTO.html>

2.4.3 Setting your PATH

The PATH environment variable helps your shell know where to look for binaries to run. It will search all directories in your \$PATH until it either finds a binary with that name or fails. You could put the following line in your .bashrc file to expand the path to include your local binaries.

```
export PATH=$PATH:~/bin
```

This will allow you to easily run programs that you put in that directory. Save your .bashrc and resource it with the following

```
$ source .bashrc
```

Make a new directory ~/bin and then copy the hello binary from earlier into it.

```
$ mkdir ~/bin
```

```
$ cp hello ~/bin
```

Now run the binary by just typing

```
$ hello
```

Notice how you didn't have to type ./ before the binary name?

2.5 Retrospective and Closing Comments

If you understand all the content that we just covered then you can now do all of your programming and file managing from the command line. Try to use the command line on your next few projects to master these skills. We have barely scratched the surface, however, and we encourage you to use the **man** pages and your favorite search engine to look into the tools and topics that we just outlined. Also note that you will learn a lot from friends, teachers, and colleagues in the years to come if you continue in the CSCI program. Good luck out there!

Bibliography

- [1] Figure 1.1 found at <http://homepage.cs.uri.edu/faculty/hamel/courses/2016/fall2016/csc301/bootcamp/sessions/session-1a.html>
- [2] Figure 2 found at <http://www.linuxtrainingacademy.com/wp-content/uploads/2014/03/linux-directory-tree.jpg>